

# Event\_Classification

July 28, 2021

```
[ ]: #-----Running the code-----#  
    """  
    Just run the script in the folder containing 'data.pickle' and 'data.csv file'  
    """  
  
    #-----OVERVIEW-----#  
    """  
    * Input data is padded with a special value to make all the sequences of equal  
    length and an additional masking layer is used to mask these special value  
    during training.  
  
    * Since the dataset is imbalanced various techniques like stratification, class  
    weighting and precision-recall curves are also used.  
  
    * Model: Stacked LSTM layers are used for encoding the time series and then the  
    encoding is passed through dense layers for classification.  
  
    * Hyperparameter tuning is performed to find the best size of the network.  
    """
```

```
[2]: import tensorflow as tf  
    from tensorflow import keras  
    from keras.layers import LSTM, Dense, Masking, BatchNormalization, Activation  
    import keras_tuner as kt  
  
    import pickle  
  
    import matplotlib.pyplot as plt  
    import numpy as np  
    import pandas as pd  
    import seaborn as sns  
  
    import sklearn  
    from sklearn.metrics import confusion_matrix  
    from sklearn.model_selection import train_test_split
```

## Helper Functions

```
[4]: def plot_loss(history):  
  
    """  
    Using the history logs of the trained model this function plots the training  
    loss and validation loss on each epoch, using a log scale on y-axis to show  
    the wide range of values.  
  
    """  
  
    # Training Loss vs Epoch  
    plt.plot(history.epoch, history.history['loss'], label='Train ')  
    # Validation Loss vs Epoch  
    plt.plot(history.epoch, history.history['val_loss'], label='Val ',  
             linestyle="--")  
  
    plt.xlabel('Epoch')  
    plt.ylabel('Loss')  
  
    plt.legend()
```

```
[39]: def plot_metrics(history):  
  
    """  
    Using the history logs and metrics of the trained model this function plots  
    4 subplots of precision, recall, loss and area under precision-recall curve  
    for the training and validation data on each epoch.  
  
    """  
  
    metrics = ['loss', 'prc', 'precision', 'recall']  
  
    for n, metric in enumerate(metrics):  
  
        plt.subplot(2,2,n+1)  
  
        plt.plot(history.epoch, history.history[metric], label='Train')  
        plt.plot(history.epoch, history.history['val_'+metric],  
                 linestyle="--", label='Val')  
  
        plt.xlabel('Epoch')  
        plt.ylabel(metric)  
  
        if metric == 'loss':  
            plt.ylim([0, plt.ylim()[1]])  
        else:  
            plt.ylim([0,1])
```

```
plt.legend()
```

```
[6]: def plot_cm(true_labels, predictions, p=0.5):  
    """  
    Using the inbuilt function of confusion_matrix from sklearn this function  
    plots a confusion matrix with a particular threshold(p) with default value  
    0.5.  
  
    """  
    # Confusion matrix  
    cm = confusion_matrix(true_labels, predictions > p)  
  
    plt.figure(figsize=(5,5))  
    sns.heatmap(cm, annot=True, fmt="d")  
    plt.title('Confusion matrix @{: .2f}'.format(p))  
    plt.ylabel('Actual label')  
    plt.xlabel('Predicted label')  
    plt.legend()  
  
    print('(True Negatives): ', cm[0][0])  
    print('(False Positives): ', cm[0][1])  
    print('(False Negatives): ', cm[1][0])  
    print('(True Positives): ', cm[1][1])  
    print('Total : ', np.sum(cm[1]))
```

```
[8]: def plot_prc(name, true_labels, predictions, **kwargs):  
    """  
    Using the inbuilt function of precision_recall_curve from sklearn this  
    function plots a PR curve  
  
    """  
    precision, recall, _ = sklearn.metrics.precision_recall_curve(true_labels,  
                                                                    predictions)  
  
    plt.plot(precision, recall, label=name, linewidth=2, **kwargs)  
    plt.xlabel('Recall')  
    plt.ylabel('Precision')  
    plt.grid(True)  
    plt.legend()  
    ax = plt.gca()  
    ax.set_aspect('equal')
```

```
[13]: def visualize_event(data, e):  
    """  
    This function visualizes all the features across timesteps in the data  
    for a particular event e.
```

```

"""

fig, axes = plt.subplots(nrows=3, ncols=3, dpi=120, figsize=(10,6))
idx = 0
for i in range(9):

    ax = axes.flatten()[idx]
    ax.plot(range(len(data[e][:,i])), data[e][:,i], linewidth=1)
    idx += 1

    ax.set_xlabel('Timesteps')
    ax.set_ylabel("Feature_" + str(i))

plt.tight_layout()

```

Building model

```

[26]: def hyper_model(hp):
    """
    This function builds a model with four hyperparameters(lstm_units,
    dense1_units, dense2_units and learning_rate). This function is
    passed to a keras hypermeter tuner to select a best hyperparameter
    combination.

    """

    model = keras.Sequential()

    # Masks a sequence by using a mask_value to skip timesteps which were padded
    model.add(Masking(mask_value=special_value,
                      input_shape=(max_seq_len, dimension)))

    # Tune the number of units in the LSTM layer
    # Choose an optimal value between 128-256
    lstm_units = hp.Int('lstm_units', min_value=128, max_value=256, step=128)
    model.add(LSTM(lstm_units, return_sequences = True))
    model.add(LSTM(lstm_units))

    # Tune the number of units in the first dense layer
    # Choose an optimal value between 256-512
    dense1_units = hp.Int('Dense_1', min_value=256, max_value=512, step=256)
    model.add(Dense(dense1_units))
    model.add(BatchNormalization())
    model.add(Activation(activation='relu'))

    # Tune the number of units in the second dense layer

```

```

# Choose an optimal value between 64-128
dense2_units = hp.Int('Dense_2', min_value=64, max_value=128, step=64)
model.add(Dense(dense2_units))
model.add(BatchNormalization())
model.add(Activation(activation='relu'))

model.add(Dense(1, activation='sigmoid'))

# Tune the learning rate for the optimizer
# Choose an optimal value from 0.001, or 0.0001
hp_learning_rate = hp.Choice('learning_rate', values=[1e-3, 1e-4])

model.compile(optimizer=keras.optimizers.Adam(
    learning_rate=hp_learning_rate),
    loss=keras.losses.BinaryCrossentropy(),
    metrics=METRICS)

return model

```

Load Data

```

[15]: # Load data.pickle as raw_data
filename = 'data.pickle'
infile = open(filename, 'rb')
raw_data = pickle.load(infile)
infile.close()

# Load data.csv as dataframe
df = pd.read_csv('data.csv')

```

Check the class distribution

```


[16]: # output label
y = np.array(df['selected_event'])

# count number of each class in the dataset
unique, counts = np.unique(y, return_counts=True)
print(f'Number of each output label are {dict(zip(unique, counts))}')

# plot the histogram for each class
plt.rcParams["figure.figsize"] = (10,6)
plt.hist(y, bins=2, density=False)
plt.ylabel('Number of samples')
plt.xlabel('(no = 0; yes = 1)')
plt.show()

```

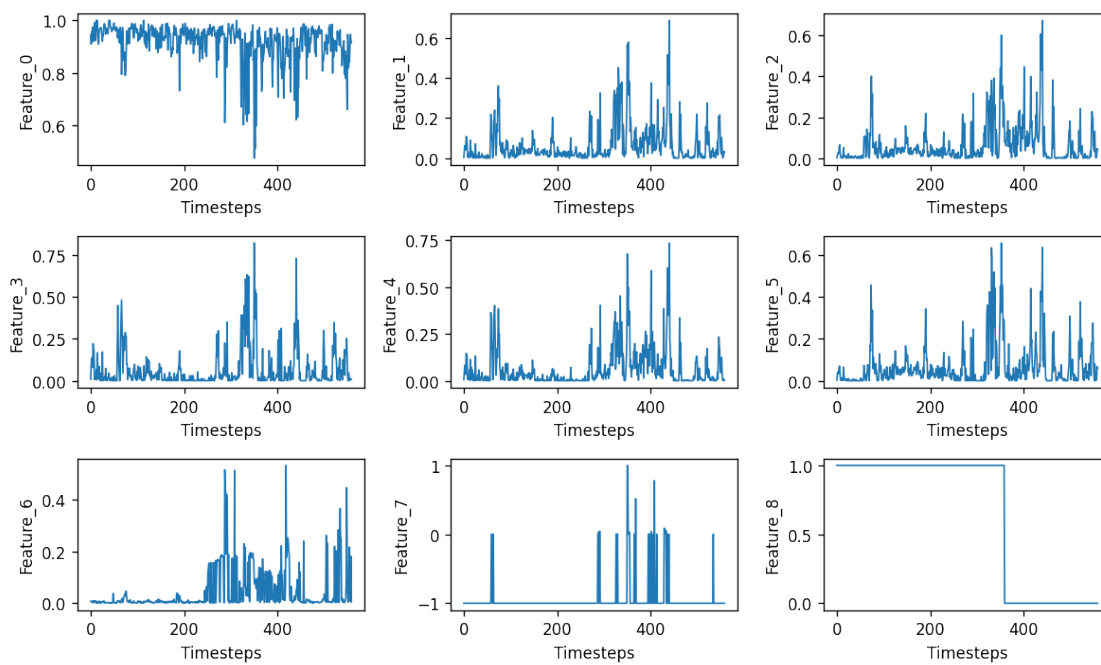
Number of each output label are {0: 1775, 1: 225}



Visualizes all the features across timesteps for a particular event

```
[18]: data = np.array(raw_data, dtype=object)

event = 100 #event to be visualaized
visualize_event(data, event)
```



Pre-process the Data

```
[19]: # Padding the data with special_value(which will be masked later in LSTM layers)
# to make all the sequences all equal length.
max_seq_len = df['X_length'].max()
special_value = -10.0

# Value 2000 and 10 are hard coded. These are number of events and feature size
# respectively.
X = np.full((2000,max_seq_len,10), special_value, dtype=float)

for i in range(data.size):
    tmp = data[i].shape
    X[i,:tmp[0],:] = data[i]

# Check for missing anf inf vlaues
X = np.nan_to_num(X, copy=True, nan=0.0, posinf=None, neginf=None)
```

Calculating class weights

```
[20]: # Since the data considered here is imbalanced, the weighted loss function is
# used where minority class is penalized more according to the class weights.
# This block calculates the class_weight which will passed to the model's loss
# function.

total = counts[0] + counts[1]
weight_for_0 = (1 / counts[0]) * (total / 2.0)
weight_for_1 = (1 / counts[1]) * (total / 2.0)

class_weight = {0: weight_for_0, 1: weight_for_1}

print('Weight for class 0: {:.2f}'.format(weight_for_0))
print('Weight for class 1: {:.2f}'.format(weight_for_1))
```

Weight for class 0: 0.56

Weight for class 1: 4.44

Train-Validation-Test split

```
[21]: # Split the dataset into train, validation and test set. Using stratify to make
# the split balanced with both classes.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1,
                                                    stratify = y)

X_train, X_val, y_train, y_val = train_test_split(X_train, y_train,
                                                    test_size=0.2,
```

```
stratify = y_train)
```

Define Parameters

```
[22]: # All the metrics which will be monitored during training.
METRICS = [
    keras.metrics.TruePositives(name='tp'),
    keras.metrics.FalsePositives(name='fp'),
    keras.metrics.TrueNegatives(name='tn'),
    keras.metrics.FalseNegatives(name='fn'),
    keras.metrics.BinaryAccuracy(name='accuracy'),
    keras.metrics.Precision(name='precision'),
    keras.metrics.Recall(name='recall'),
    keras.metrics.AUC(name='auc'),
    keras.metrics.AUC(name='prc', curve='PR'), # precision-recall curve
]
```

```
[23]: # Parameters
dimension = 10
verbosity = 1
EPOCHS = 200
BATCH_SIZE = 256

# Callbacks used during training
# Validation precision-recall curve is monitored instead of accuracy since
# the data is imbalanced
early_stopping = tf.keras.callbacks.EarlyStopping(monitor='val_prc',
                                                    verbose=1,
                                                    patience=10, mode='max')

# Callback for saving the best model
chk = tf.keras.callbacks.ModelCheckpoint('best_model',
                                         monitor='val_prc',
                                         save_best_only=True,
                                         mode='max', verbose=1)
```

Hyperparameter Tuning using Bayesian Optimization

```
[27]: # Hyperparameter tuning using Bayesian Optimization for hyper_model
# defined above
tuner = kt.BayesianOptimization(hyper_model, max_trials=10,
                                objective=kt.Objective("val_prc",
                                                        direction="max"))

# Hyperparameter search according to the BayesianOptimization tuner
tuner.search(X_train, y_train, epochs=10, batch_size=BATCH_SIZE,
             class_weight=class_weight,
```



```

        validation_data=(X_val, y_val), callbacks=[early_stopping])

# Get the optimal hyperparameters
best_hps = tuner.get_best_hyperparameters(num_trials=1)[0]

print(f"""
The hyperparameter search is complete. The optimal number of units
in the lstm output is {best_hps.get('lstm_units')}, first densely-connected
layer is {best_hps.get('Dense_1')} and second densely-connected layer is
{best_hps.get('Dense_2')} and the optimal learning rate for the optimizer
is {best_hps.get('learning_rate')}.
""")

```

Trial 10 Complete [00h 22m 33s]  
val\_prc: 0.17031022906303406

Best val\_prc So Far: 0.8057265877723694  
Total elapsed time: 03h 14m 04s  
INFO:tensorflow:Oracle triggered exit

The hyperparameter search is complete. The optimal number of units in the lstm output is 256, first densely-connected layer is 256 and second densely-connected layer is 128 and the optimal learning rate for the optimizer is 0.001.

## Model Training

```

[33]: # Build the model with the optimal hyperparameters and train it on the data
model = tuner.hypermodel.build(best_hps)
history = model.fit(X_train, y_train, batch_size=BATCH_SIZE, epochs=100,
                    validation_data=(X_val, y_val),
                    class_weight=class_weight, callbacks=[early_stopping, chk])

```

Epoch 1/100

6/6 [=====] - 154s 24s/step - loss: 0.5346 - tp: 153.0000 - fp: 708.0000 - tn: 889.0000 - fn: 50.0000 - accuracy: 0.5789 - precision: 0.1777 - recall: 0.7537 - auc: 0.7574 - prc: 0.4443 - val\_loss: 0.6110 - val\_tp: 1.0000 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 40.0000 - val\_accuracy: 0.8889 - val\_precision: 1.0000 - val\_recall: 0.0244 - val\_auc: 0.7979 - val\_prc: 0.6883

Epoch 00001: val\_prc improved from -inf to 0.68828, saving model to /content/drive/MyDrive/Experiments/Experiment\_4/best\_model

WARNING:absl:Found untraced functions such as masking\_layer\_call\_and\_return\_conditional\_losses, masking\_layer\_call\_fn, dense\_layer\_call\_and\_return\_conditional\_losses, dense\_layer\_call\_fn,

activation\_layer\_call\_and\_return\_conditional\_losses while saving (showing 5 of 40). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to:

/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

INFO:tensorflow:Assets written to:

/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

Epoch 2/100

6/6 [=====] - 147s 24s/step - loss: 0.3735 - tp: 138.0000 - fp: 207.0000 - tn: 1071.0000 - fn: 24.0000 - accuracy: 0.8396 - precision: 0.4000 - recall: 0.8519 - auc: 0.9261 - prc: 0.7045 - val\_loss: 0.4228 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.8840 - val\_prc: 0.5224

Epoch 00002: val\_prc did not improve from 0.68828

Epoch 3/100

6/6 [=====] - 154s 25s/step - loss: 0.3806 - tp: 121.0000 - fp: 141.0000 - tn: 1137.0000 - fn: 41.0000 - accuracy: 0.8736 - precision: 0.4618 - recall: 0.7469 - auc: 0.9055 - prc: 0.7155 - val\_loss: 0.4170 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.8801 - val\_prc: 0.6697

Epoch 00003: val\_prc did not improve from 0.68828

Epoch 4/100

6/6 [=====] - 157s 26s/step - loss: 0.3549 - tp: 137.0000 - fp: 160.0000 - tn: 1118.0000 - fn: 25.0000 - accuracy: 0.8715 - precision: 0.4613 - recall: 0.8457 - auc: 0.9267 - prc: 0.7267 - val\_loss: 0.4263 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.8024 - val\_prc: 0.6616

Epoch 00004: val\_prc did not improve from 0.68828

Epoch 5/100

6/6 [=====] - 149s 25s/step - loss: 0.3303 - tp: 139.0000 - fp: 161.0000 - tn: 1117.0000 - fn: 23.0000 - accuracy: 0.8722 - precision: 0.4633 - recall: 0.8580 - auc: 0.9343 - prc: 0.7544 - val\_loss: 0.3639 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.7900 - val\_prc: 0.6627

Epoch 00005: val\_prc did not improve from 0.68828

Epoch 6/100

6/6 [=====] - 152s 25s/step - loss: 0.3062 - tp: 139.0000 - fp: 161.0000 - tn: 1117.0000 - fn: 23.0000 - accuracy: 0.8722 - precision: 0.4633 - recall: 0.8580 - auc: 0.9468 - prc: 0.7568 - val\_loss:

0.3381 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn:  
41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall:  
0.0000e+00 - val\_auc: 0.8459 - val\_prc: 0.7266

Epoch 00006: val\_prc improved from 0.68828 to 0.72660, saving model to  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model

WARNING:absl:Found untraced functions such as  
masking\_layer\_call\_and\_return\_conditional\_losses, masking\_layer\_call\_fn,  
dense\_layer\_call\_and\_return\_conditional\_losses, dense\_layer\_call\_fn,  
activation\_layer\_call\_and\_return\_conditional\_losses while saving (showing 5 of  
40). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

Epoch 7/100

6/6 [=====] - 147s 25s/step - loss: 0.2880 - tp:  
136.0000 - fp: 121.0000 - tn: 1157.0000 - fn: 26.0000 - accuracy: 0.8979 -  
precision: 0.5292 - recall: 0.8395 - auc: 0.9501 - prc: 0.8017 - val\_loss:  
0.3295 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn:  
41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall:  
0.0000e+00 - val\_auc: 0.8487 - val\_prc: 0.7309

Epoch 00007: val\_prc improved from 0.72660 to 0.73085, saving model to  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model

WARNING:absl:Found untraced functions such as  
masking\_layer\_call\_and\_return\_conditional\_losses, masking\_layer\_call\_fn,  
dense\_layer\_call\_and\_return\_conditional\_losses, dense\_layer\_call\_fn,  
activation\_layer\_call\_and\_return\_conditional\_losses while saving (showing 5 of  
40). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

Epoch 8/100

6/6 [=====] - 139s 23s/step - loss: 0.2755 - tp:  
143.0000 - fp: 121.0000 - tn: 1157.0000 - fn: 19.0000 - accuracy: 0.9028 -  
precision: 0.5417 - recall: 0.8827 - auc: 0.9557 - prc: 0.7921 - val\_loss:  
0.3621 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn:  
41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall:  
0.0000e+00 - val\_auc: 0.7918 - val\_prc: 0.6727

Epoch 00008: val\_prc did not improve from 0.73085

Epoch 9/100

6/6 [=====] - 147s 25s/step - loss: 0.2673 - tp: 141.0000 - fp: 118.0000 - tn: 1160.0000 - fn: 21.0000 - accuracy: 0.9035 - precision: 0.5444 - recall: 0.8704 - auc: 0.9595 - prc: 0.7937 - val\_loss: 0.4120 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.8579 - val\_prc: 0.6128

Epoch 00009: val\_prc did not improve from 0.73085

Epoch 10/100

6/6 [=====] - 144s 24s/step - loss: 0.3174 - tp: 136.0000 - fp: 140.0000 - tn: 1138.0000 - fn: 26.0000 - accuracy: 0.8847 - precision: 0.4928 - recall: 0.8395 - auc: 0.9386 - prc: 0.7585 - val\_loss: 0.3547 - val\_tp: 7.0000 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 34.0000 - val\_accuracy: 0.9056 - val\_precision: 1.0000 - val\_recall: 0.1707 - val\_auc: 0.7860 - val\_prc: 0.6359

Epoch 00010: val\_prc did not improve from 0.73085

Epoch 11/100

6/6 [=====] - 138s 23s/step - loss: 0.3012 - tp: 136.0000 - fp: 147.0000 - tn: 1131.0000 - fn: 26.0000 - accuracy: 0.8799 - precision: 0.4806 - recall: 0.8395 - auc: 0.9464 - prc: 0.7660 - val\_loss: 0.2849 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.9195 - val\_prc: 0.6510

Epoch 00011: val\_prc did not improve from 0.73085

Epoch 12/100

6/6 [=====] - 139s 23s/step - loss: 0.2883 - tp: 136.0000 - fp: 105.0000 - tn: 1173.0000 - fn: 26.0000 - accuracy: 0.9090 - precision: 0.5643 - recall: 0.8395 - auc: 0.9540 - prc: 0.7734 - val\_loss: 0.3802 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.8798 - val\_prc: 0.6111

Epoch 00012: val\_prc did not improve from 0.73085

Epoch 13/100

6/6 [=====] - 141s 24s/step - loss: 0.3003 - tp: 146.0000 - fp: 217.0000 - tn: 1061.0000 - fn: 16.0000 - accuracy: 0.8382 - precision: 0.4022 - recall: 0.9012 - auc: 0.9497 - prc: 0.7938 - val\_loss: 0.3219 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.9125 - val\_prc: 0.7526

Epoch 00013: val\_prc improved from 0.73085 to 0.75259, saving model to /content/drive/MyDrive/Experiments/Experiment\_4/best\_model

WARNING:absl:Found untraced functions such as masking\_layer\_call\_and\_return\_conditional\_losses, masking\_layer\_call\_fn,

dense\_layer\_call\_and\_return\_conditional\_losses, dense\_layer\_call\_fn, activation\_layer\_call\_and\_return\_conditional\_losses while saving (showing 5 of 40). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to:

/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

INFO:tensorflow:Assets written to:

/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

Epoch 14/100

6/6 [=====] - 138s 23s/step - loss: 0.2883 - tp: 139.0000 - fp: 98.0000 - tn: 1180.0000 - fn: 23.0000 - accuracy: 0.9160 - precision: 0.5865 - recall: 0.8580 - auc: 0.9495 - prc: 0.7748 - val\_loss: 0.3293 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.7735 - val\_prc: 0.4724

Epoch 00014: val\_prc did not improve from 0.75259

Epoch 15/100

6/6 [=====] - 138s 23s/step - loss: 0.2927 - tp: 141.0000 - fp: 150.0000 - tn: 1128.0000 - fn: 21.0000 - accuracy: 0.8813 - precision: 0.4845 - recall: 0.8704 - auc: 0.9476 - prc: 0.7858 - val\_loss: 0.2809 - val\_tp: 5.0000 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 36.0000 - val\_accuracy: 0.9000 - val\_precision: 1.0000 - val\_recall: 0.1220 - val\_auc: 0.9267 - val\_prc: 0.7534

Epoch 00015: val\_prc improved from 0.75259 to 0.75336, saving model to

/content/drive/MyDrive/Experiments/Experiment\_4/best\_model

WARNING:absl:Found untraced functions such as

masking\_layer\_call\_and\_return\_conditional\_losses, masking\_layer\_call\_fn, dense\_layer\_call\_and\_return\_conditional\_losses, dense\_layer\_call\_fn, activation\_layer\_call\_and\_return\_conditional\_losses while saving (showing 5 of 40). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to:

/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

INFO:tensorflow:Assets written to:

/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

Epoch 16/100

6/6 [=====] - 136s 23s/step - loss: 0.2657 - tp: 143.0000 - fp: 132.0000 - tn: 1146.0000 - fn: 19.0000 - accuracy: 0.8951 - precision: 0.5200 - recall: 0.8827 - auc: 0.9582 - prc: 0.8119 - val\_loss: 0.2928 - val\_tp: 11.0000 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 30.0000 - val\_accuracy: 0.9167 - val\_precision: 1.0000 - val\_recall: 0.2683 - val\_auc: 0.9485 - val\_prc: 0.8182

Epoch 00016: val\_prc improved from 0.75336 to 0.81823, saving model to

/content/drive/MyDrive/Experiments/Experiment\_4/best\_model

WARNING:absl:Found untraced functions such as  
masking\_layer\_call\_and\_return\_conditional\_losses, masking\_layer\_call\_fn,  
dense\_layer\_call\_and\_return\_conditional\_losses, dense\_layer\_call\_fn,  
activation\_layer\_call\_and\_return\_conditional\_losses while saving (showing 5 of  
40). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

Epoch 17/100

6/6 [=====] - 143s 24s/step - loss: 0.2822 - tp:  
141.0000 - fp: 121.0000 - tn: 1157.0000 - fn: 21.0000 - accuracy: 0.9014 -  
precision: 0.5382 - recall: 0.8704 - auc: 0.9502 - prc: 0.8049 - val\_loss:  
0.2863 - val\_tp: 22.0000 - val\_fp: 1.0000 - val\_tn: 318.0000 - val\_fn: 19.0000 -  
val\_accuracy: 0.9444 - val\_precision: 0.9565 - val\_recall: 0.5366 - val\_auc:  
0.9539 - val\_prc: 0.8336

Epoch 00017: val\_prc improved from 0.81823 to 0.83359, saving model to  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model

WARNING:absl:Found untraced functions such as  
masking\_layer\_call\_and\_return\_conditional\_losses, masking\_layer\_call\_fn,  
dense\_layer\_call\_and\_return\_conditional\_losses, dense\_layer\_call\_fn,  
activation\_layer\_call\_and\_return\_conditional\_losses while saving (showing 5 of  
40). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

Epoch 18/100

6/6 [=====] - 134s 22s/step - loss: 0.2785 - tp:  
137.0000 - fp: 107.0000 - tn: 1171.0000 - fn: 25.0000 - accuracy: 0.9083 -  
precision: 0.5615 - recall: 0.8457 - auc: 0.9536 - prc: 0.7945 - val\_loss:  
0.2871 - val\_tp: 6.0000 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn:  
35.0000 - val\_accuracy: 0.9028 - val\_precision: 1.0000 - val\_recall: 0.1463 -  
val\_auc: 0.9354 - val\_prc: 0.7713

Epoch 00018: val\_prc did not improve from 0.83359

Epoch 19/100

6/6 [=====] - 135s 23s/step - loss: 0.2881 - tp:  
142.0000 - fp: 172.0000 - tn: 1106.0000 - fn: 20.0000 - accuracy: 0.8667 -  
precision: 0.4522 - recall: 0.8765 - auc: 0.9510 - prc: 0.7864 - val\_loss:  
0.3325 - val\_tp: 11.0000 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn:  
30.0000 - val\_accuracy: 0.9167 - val\_precision: 1.0000 - val\_recall: 0.2683 -

val\_auc: 0.9574 - val\_prc: 0.8448

Epoch 00019: val\_prc improved from 0.83359 to 0.84482, saving model to  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model

WARNING:absl:Found untraced functions such as  
masking\_layer\_call\_and\_return\_conditional\_losses, masking\_layer\_call\_fn,  
dense\_layer\_call\_and\_return\_conditional\_losses, dense\_layer\_call\_fn,  
activation\_layer\_call\_and\_return\_conditional\_losses while saving (showing 5 of  
40). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

INFO:tensorflow:Assets written to:  
/content/drive/MyDrive/Experiments/Experiment\_4/best\_model/assets

Epoch 20/100

6/6 [=====] - 129s 22s/step - loss: 0.2759 - tp:  
144.0000 - fp: 182.0000 - tn: 1096.0000 - fn: 18.0000 - accuracy: 0.8611 -  
precision: 0.4417 - recall: 0.8889 - auc: 0.9549 - prc: 0.8107 - val\_loss:  
0.3306 - val\_tp: 3.0000 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn:  
38.0000 - val\_accuracy: 0.8944 - val\_precision: 1.0000 - val\_recall: 0.0732 -  
val\_auc: 0.9530 - val\_prc: 0.8198

Epoch 00020: val\_prc did not improve from 0.84482

Epoch 21/100

6/6 [=====] - 138s 23s/step - loss: 0.2521 - tp:  
145.0000 - fp: 125.0000 - tn: 1153.0000 - fn: 17.0000 - accuracy: 0.9014 -  
precision: 0.5370 - recall: 0.8951 - auc: 0.9631 - prc: 0.8105 - val\_loss:  
0.2070 - val\_tp: 19.0000 - val\_fp: 1.0000 - val\_tn: 318.0000 - val\_fn: 22.0000 -  
val\_accuracy: 0.9361 - val\_precision: 0.9500 - val\_recall: 0.4634 - val\_auc:  
0.9562 - val\_prc: 0.8350

Epoch 00021: val\_prc did not improve from 0.84482

Epoch 22/100

6/6 [=====] - 130s 22s/step - loss: 0.3083 - tp:  
132.0000 - fp: 86.0000 - tn: 1192.0000 - fn: 30.0000 - accuracy: 0.9194 -  
precision: 0.6055 - recall: 0.8148 - auc: 0.9507 - prc: 0.7680 - val\_loss:  
0.3581 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn:  
41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall:  
0.0000e+00 - val\_auc: 0.8267 - val\_prc: 0.6374

Epoch 00022: val\_prc did not improve from 0.84482

Epoch 23/100

6/6 [=====] - 135s 22s/step - loss: 0.3071 - tp:  
130.0000 - fp: 110.0000 - tn: 1168.0000 - fn: 32.0000 - accuracy: 0.9014 -  
precision: 0.5417 - recall: 0.8025 - auc: 0.9425 - prc: 0.7820 - val\_loss:  
0.3154 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn:  
41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall:

0.0000e+00 - val\_auc: 0.8229 - val\_prc: 0.4491

Epoch 00023: val\_prc did not improve from 0.84482

Epoch 24/100

6/6 [=====] - 131s 22s/step - loss: 0.2925 - tp: 144.0000 - fp: 170.0000 - tn: 1108.0000 - fn: 18.0000 - accuracy: 0.8694 - precision: 0.4586 - recall: 0.8889 - auc: 0.9481 - prc: 0.7946 - val\_loss: 0.3199 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.8235 - val\_prc: 0.3568

Epoch 00024: val\_prc did not improve from 0.84482

Epoch 25/100

6/6 [=====] - 136s 23s/step - loss: 0.2752 - tp: 147.0000 - fp: 181.0000 - tn: 1097.0000 - fn: 15.0000 - accuracy: 0.8639 - precision: 0.4482 - recall: 0.9074 - auc: 0.9559 - prc: 0.8171 - val\_loss: 0.2971 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.8571 - val\_prc: 0.4360

Epoch 00025: val\_prc did not improve from 0.84482

Epoch 26/100

6/6 [=====] - 144s 24s/step - loss: 0.2732 - tp: 140.0000 - fp: 108.0000 - tn: 1170.0000 - fn: 22.0000 - accuracy: 0.9097 - precision: 0.5645 - recall: 0.8642 - auc: 0.9543 - prc: 0.7969 - val\_loss: 0.3008 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.8659 - val\_prc: 0.4358

Epoch 00026: val\_prc did not improve from 0.84482

Epoch 27/100

6/6 [=====] - 133s 22s/step - loss: 0.2906 - tp: 135.0000 - fp: 120.0000 - tn: 1158.0000 - fn: 27.0000 - accuracy: 0.8979 - precision: 0.5294 - recall: 0.8333 - auc: 0.9483 - prc: 0.7797 - val\_loss: 0.3196 - val\_tp: 0.0000e+00 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 41.0000 - val\_accuracy: 0.8861 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00 - val\_auc: 0.8409 - val\_prc: 0.5445

Epoch 00027: val\_prc did not improve from 0.84482

Epoch 28/100

6/6 [=====] - 135s 22s/step - loss: 0.2756 - tp: 139.0000 - fp: 115.0000 - tn: 1163.0000 - fn: 23.0000 - accuracy: 0.9042 - precision: 0.5472 - recall: 0.8580 - auc: 0.9554 - prc: 0.7997 - val\_loss: 0.2939 - val\_tp: 9.0000 - val\_fp: 1.0000 - val\_tn: 318.0000 - val\_fn: 32.0000 - val\_accuracy: 0.9083 - val\_precision: 0.9000 - val\_recall: 0.2195 - val\_auc: 0.9257 - val\_prc: 0.6689

Epoch 00028: val\_prc did not improve from 0.84482




Epoch 29/100

6/6 [=====] - 139s 23s/step - loss: 0.2611 - tp: 145.0000 - fp: 135.0000 - tn: 1143.0000 - fn: 17.0000 - accuracy: 0.8944 - precision: 0.5179 - recall: 0.8951 - auc: 0.9585 - prc: 0.8156 - val\_loss: 0.2738 - val\_tp: 3.0000 - val\_fp: 0.0000e+00 - val\_tn: 319.0000 - val\_fn: 38.0000 - val\_accuracy: 0.8944 - val\_precision: 1.0000 - val\_recall: 0.0732 - val\_auc: 0.8595 - val\_prc: 0.7224

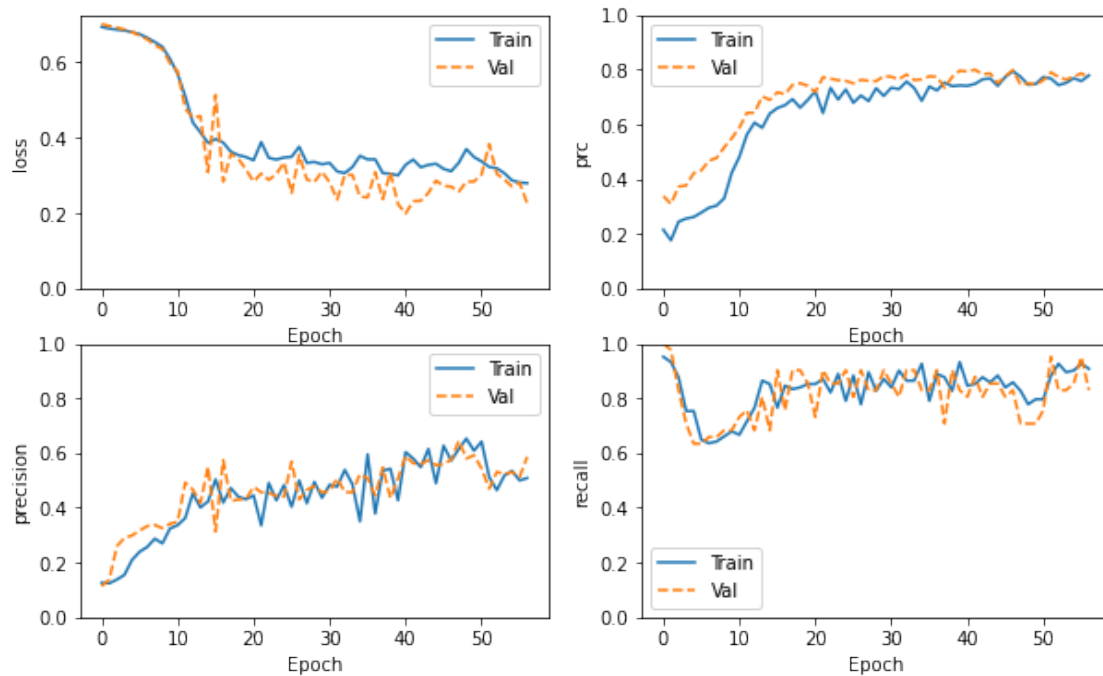
Epoch 00029: val\_prc did not improve from 0.84482

Epoch 00029: early stopping

```
[47]: # Loss plot  
plot_loss(history)
```



```
[48]: # Metric plot  
plot_metrics(history)
```




Evaluate Model

```
[41]: # Predict for train and test data using trained model
train_predictions = model.predict(X_train, batch_size=BATCH_SIZE)
test_predictions = model.predict(X_test, batch_size=BATCH_SIZE)
```

```
[51]: # Plot Confusion Matrix
plot_cm(y_test, test_predictions)
```

WARNING:matplotlib.legend:No handles with labels found to put in legend.

```
(True Negatives): 166
(False Positives): 12
(False Negatives): 1
(True Positives): 21
Total : 22
```



```
[50]: # Evaluate the model and output the various metrics
results = model.evaluate(X_test, y_test, batch_size=BATCH_SIZE, verbose=0)

for name, value in zip(model.metrics_names, results):
    print(name, ': ', value)
print()
```

```
loss : 0.18149976432323456
tp : 21.0
fp : 12.0
tn : 166.0
fn : 1.0
accuracy : 0.9350000023841858
precision : 0.6363636255264282
recall : 0.9545454382896423
auc : 0.9802093505859375
prc : 0.8792988061904907
```

```
[45]: # Plot Precision-recall curve
plot_prc("Train", y_train, train_predictions)
plot_prc("Test", y_test, test_predictions, linestyle='--')
```

```
[45]: <matplotlib.legend.Legend at 0x7f940fec9b10>
```

